

Designs with two factors

Last time: Completely randomized designs

- ▶ Two equivalent parametrizations:

1. Group means parametrization:

$$Y_{ij} = \mu_i + \varepsilon_{ij}, \quad \varepsilon_{ij} \stackrel{\text{iid}}{\sim} N(0, \sigma^2),$$

2. Sum-to-zero parametrization:

$$Y_{ij} = \mu + \tau_i + \varepsilon_{ij}, \quad \sum_{i=1}^a \tau_i = 0, \quad \varepsilon_{ij} \stackrel{\text{iid}}{\sim} N(0, \sigma^2),$$

- ▶ $i \in \{1, \dots, a\}$ indexes group membership
- ▶ $j \in \{1, \dots, n_i\}$ indexes observations within a group

Today

Designs with two factors

- ▶ **One treatment and a block**
- ▶ **Two treatments**

One treatment and a block

One treatment and a block

- ▶ We want to quantify the effect of the treatment on the response
- ▶ In an experimental design, we can control the assignment mechanism of the treatment
- ▶ A block (or blocking variable) is a known source of variation we can control (assign). It affects the outcome but we are not interested in its effect

Example: Painting line

- ▶ An engineer working at a painting line wants to know if different types of paints have the same average drying times
- ▶ They will perform the experiments on four different days
- ▶ The engineers think that the weather may affect the outcome: that is, the day may affect the outcome
- ▶ The treatment is the type of paint, the block is the day

Balanced complete randomized block designs with replicates

- ▶ We assign experimental units randomly to blocks and treatments
- ▶ The block has b levels
- ▶ The treatment has a levels
- ▶ We collect r observations for each combination of block and treatment: there are r replicates
- ▶ Total sample size is $N = rba$

Terminology

- ▶ The design is *balanced* because each combination of block and treatment is observed the same number of times: r
- ▶ The design is *complete* because there are observations for each combination of block and treatment
- ▶ The term “balanced complete randomized block designs” usually refers to the case $r = 1$, which is the most common
- ▶ Designs with $r > 1$ are often referred to as “balanced complete randomized block designs with replicates”

Example: Painting line

- ▶ Day (block): $b = 4$ days
- ▶ Type of paint (treatment): $a = 3$ types
- ▶ $r = 1$ replicate
- ▶ $N = rba = 12$ observations

	Day 1	Day 2	Day 3	Day 4
Type A	2.60	2.10	3.5	2.65
Type B	3.25	2.65	3.8	3.50
Type C	2.60	2.05	3.1	2.62

Model

$$Y_{ijk} = \mu + \tau_i + \beta_j + \varepsilon_{ijk}, \quad \sum_{i=1}^a \tau_i = \sum_{j=1}^b \beta_j = 0, \quad \varepsilon_{ijk} \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$$

- ▶ $i \in \{1, \dots, a\}$ indexes treatments, $j \in \{1, \dots, b\}$ indexes blocks, and $k \in \{1, \dots, r\}$ indexes replicates
- ▶ As before, we assume that the errors are independent and normal, and that the variance σ^2 doesn't depend on the treatment or block
- ▶ We also assume an additive relationship between treatment and block

Additivity

- ▶ Treatment effect doesn't depend on blocks; block effect doesn't depend on treatment
- ▶ **Example:** $a = 2$ treatments and $b = 2$ blocks, table with expected values

	Block 1	Block 2
Treat 1	$\mu + \tau_1 + \beta_1$	$\mu + \tau_1 + \beta_1$
Treat 2	$\mu + \tau_2 + \beta_2$	$\mu + \tau_2 + \beta_2$


Average treatment effect in block 1: $\mathbb{E}(y_{11k} - y_{21k'}) = \tau_1 - \tau_2$

Average treatment effect in block 2: $\mathbb{E}(y_{12k} - y_{22k'}) = \tau_1 - \tau_2$

Average block effect in treat 1: $\mathbb{E}(y_{11k} - y_{12k'}) = \beta_1 - \beta_2$

Average block effect in treat 2: $\mathbb{E}(y_{21k} - y_{22k'}) = \beta_1 - \beta_2$

Example: estimated effects with additive model



Estimating parameters: Block design

► Let

$$\text{total mean} = \bar{Y} = \frac{1}{N} \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r Y_{ijk}$$

$$\text{treat } i \text{ mean} = \bar{Y}_{i.} = \frac{1}{br} \sum_{j=1}^b \sum_{k=1}^r Y_{ijk}$$

$$\text{block } j \text{ mean} = \bar{Y}_{.j} = \frac{1}{ar} \sum_{i=1}^a \sum_{k=1}^r Y_{ijk}$$

► Point estimates

$$\hat{\mu} = \bar{Y}, \quad \hat{\tau}_i = \bar{Y}_{i.} - \bar{Y}, \quad \hat{\beta}_j = \bar{Y}_{.j} - \bar{Y},$$

and

$$\hat{\sigma}^2 = \frac{1}{N-1} \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (Y_{ijk} - \hat{\mu} - \hat{\tau}_i - \hat{\beta}_j)^2$$

Example: Reading in data

```
mat = matrix(c(2.60, 2.10, 3.50, 2.65,  
              3.25, 2.65, 3.80, 3.5,  
              2.60, 2.05, 3.10, 2.62),  
            nrow = 3,  
            ncol = 4,  
            byrow = TRUE)  
rownames(mat) = c("Type A", "Type B", "Type C")  
colnames(mat) = c("Day 1", "Day 2", "Day 3", "Day 4")
```

Example: Converting into convenient (long) format

```
library(tidyverse)
df = as.data.frame(mat)
df$type = rownames(df)
df = df %>% pivot_longer(!type,
                        names_to = "day",
                        values_to = "outcome")
df$type = factor(df$type); df$day = factor(df$day)
```

Example: model with sum-to-zero constraint

```
# sum-to-zero
options(contrasts = c("contr.sum", "contr.poly"))
# fit model
mod = aov(outcome ~ day + type, data = df)
# point estimates
dummy.coef(mod)
```

Full coefficients are

##

(Intercept): 2.868333

day: Day 1 Day 2 Day 3 Day 4

-0.05166667 -0.60166667 0.59833333 0.05500000

type: Type A Type B Type C

-0.1558333 0.4316667 -0.2758333


Example: 95% confidence intervals

##	2.5 %	97.5 %
## (Intercept)	2.7760320	2.96063468
## day1	-0.2115373	0.10820396
## day2	-0.7615373	-0.44179604
## day3	0.4384627	0.75820396
## type1	-0.2863672	-0.02529951
## type2	0.3011328	0.56220049

No intervals for $\tau_a = -\sum_{i=1}^{a-1} \tau_i$ and $\beta_b = -\sum_{j=1}^{b-1} \beta_j$

Effect plots

```
library(emmeans)  
emmip(mod, day ~ type)
```



Global F tests

Now we have two F tests

- ▶ One for τ_i

$$H_{0,\tau} : \tau_i = 0, \text{ for all } i \in \{1, \dots, a\}$$

$$H_{1,\tau} : \text{at least one } \tau_i \neq 0$$

- ▶ Another one for β_j :

$$H_{0,\beta} : \beta_j = 0, \text{ for all } j \in \{1, \dots, b\}$$

$$H_{1,\beta} : \text{at least one } \beta_j \neq 0.$$

Sums of squares

$$SS_{\text{total}} = SS_{\text{treat}} + SS_{\text{block}} + SS_{\text{error}},$$

where

$$SS_{\text{total}} = \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (\bar{Y}_{.j} - \bar{Y})^2$$

$$SS_{\text{treat}} = br \sum_{i=1}^a (\bar{Y}_{i.} - \bar{Y})^2$$

$$SS_{\text{block}} = ar \sum_{j=1}^b (\bar{Y}_{.j} - \bar{Y})^2$$

$$SS_{\text{error}} = \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (Y_{ijk} - \hat{\mu} - \hat{\tau}_i - \hat{\beta}_j)^2$$

Under $H_{0,\tau}$, $\bar{Y}_{i.} \approx \bar{Y}$ and SS_{treat} will be small; same story with $H_{0,\beta}$ and SS_{block}

ANOVA table

There are two F statistics: one for the treatment and another one for the block

	df	SS	MS	F	p
Treat	$a - 1$	SS_{treat}	$MS_{\text{treat}} = \frac{SS_{\text{treat}}}{a-1}$	$f_{\tau,\text{obs}} = \frac{MS_{\text{treat}}}{MS_{\text{error}}}$	p_{τ}
Block	$b - 1$	SS_{block}	$MS_{\text{block}} = \frac{SS_{\text{block}}}{b-1}$	$f_{\beta,\text{obs}} = \frac{MS_{\text{block}}}{MS_{\text{error}}}$	p_{β}
Error	ν	SS_{error}	$MS_{\text{error}} = \frac{SS_{\text{error}}}{\nu}$		

where $\nu = N - a - b + 1$,

$$p_{\tau} = P(F_{a-1,\nu} > f_{\tau,\text{obs}}), \quad p_{\beta} = P(F_{b-1,\nu} > f_{\beta,\text{obs}}),$$

and $F_{\alpha,\nu}$ is notation for the F -distribution with α and ν degrees of freedom.

Example: ANOVA table

```
summary(mod)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## day           3  2.1771   0.7257   42.50 0.000195 ***
## type          2  1.1468   0.5734   33.58 0.000552 ***
## Residuals     6  0.1024   0.0171
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

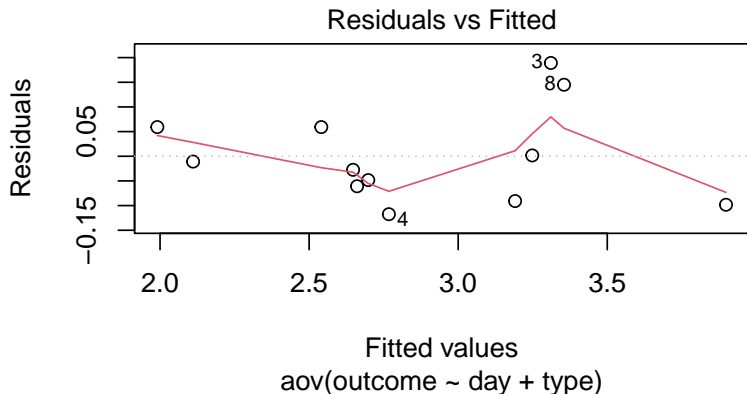
Block and treatment effect are significant at the $\alpha = 0.05$ significance level

Checking model assumptions

- ▶ We can check the hypothesis of equality of variances by looking at a plot of fitted values vs residuals: if the assumption is satisfied, the variance on the y -axis shouldn't depend much on where we are on the x -axis
- ▶ As before, we can check normality by looking at a qq-plot of the residuals
- ▶ Checking independence is hard unless we have a variable that allows us to check for space- or time-dependence, for example
- ▶ We assume that the block effect is additive. If $r = 1$, this assumption can't be checked. If $r > 1$, we can check whether an interaction term is needed or not.

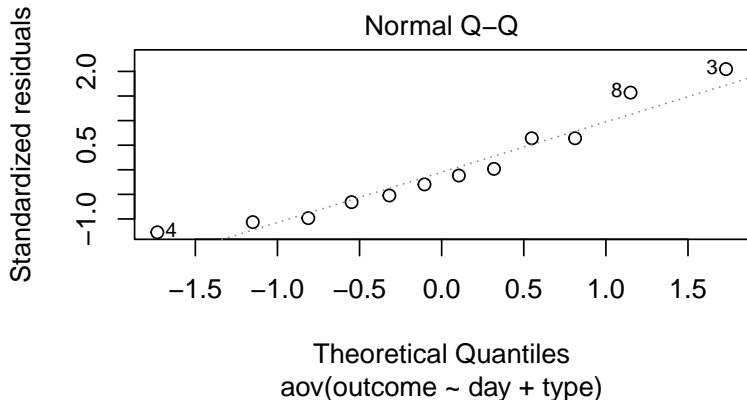
Example: checking equality of variances

```
plot(mod, which = 1)
```



Example: checking normality of residuals

```
plot(mod, which = 2)
```



Pairwise comparisons

- ▶ The global F test tells us that there are differences between treatments, but doesn't tell us where the differences are
- ▶ We can do pairwise tests to compare treatments; however, we have to be careful and control the type I error rate appropriately
- ▶ We can do that with TukeyHSD, as we did with completely randomized designs

Example: TukeyHSD

```
TukeyHSD(mod, which = "type")
```

```
## Tukey multiple comparisons of means
```

```
## 95% family-wise confidence level
```

```
##
```

```
## Fit: aov(formula = outcome ~ day + type, data = df)
```

```
##
```

```
## $type
```

##		diff	lwr	upr	p adj
## Type B-Type A	0.5875	0.3039958	0.8710042	0.0017250	
## Type C-Type A	-0.1200	-0.4035042	0.1635042	0.4459693	
## Type C-Type B	-0.7075	-0.9910042	-0.4239958	0.0006345	

Significant differences between types A and B and types B and C at the $\alpha = 0.05$ significance level

Two treatments

Two treatments


- ▶ Now we have two treatments we can control
- ▶ We are interested in their effect on the outcome
- ▶ The treatments may **interact**: the effect of a treatment may depend on the levels of the other treatment
- ▶ To be able to estimate the interaction, we need to have more than one observation for each combination of two treatments

Example: Batteries (Montgomery)

- ▶ A company is producing batteries that need to work at different temperatures and for different types of devices
- ▶ Battery life is suspected to depend on the temperature and the material that the device is made of (type A, B, or C)
- ▶ The company wants to know how temperature and the device affect battery life; they also want to know if there is an interaction between temperature and type of device
- ▶ If there is an interaction, it means that the effect of temperature depends on the type of device

Plotting the data

- There are 36 observations: 3 replicates for each combination of type and temperature



Model

$$Y_{ijk} = \mu + \tau_i + \gamma_j + (\tau\gamma)_{ij} + \varepsilon_{ijk}, \quad \varepsilon_{ijk} \stackrel{\text{iid}}{\sim} N(0, \sigma^2),$$

with sum-to-zero constraints

$$\sum_{i=1}^a \tau_i = \sum_{j=1}^b \gamma_j = 0, \quad \sum_{i=1}^a (\tau\gamma)_{ij} = \sum_{j=1}^b (\tau\gamma)_{ij} = 0$$

- ▶ $i \in \{1, \dots, a\}$ indexes the first treatment, $j \in \{1, \dots, b\}$ indexes the second treatment, and $k \in \{1, \dots, r\}$ indexes replicates
- ▶ The errors are independent and normal, the variance σ^2 doesn't depend on the treatment
- ▶ $(\tau\gamma)_{ij}$ is the interaction term

Point estimation

Let

$$\begin{aligned}\bar{Y} &= \frac{1}{N} \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r Y_{ijk}, & \bar{Y}_{ij} &= \frac{1}{r} \sum_{k=1}^r Y_{ijk} \\ \bar{Y}_{i.} &= \frac{1}{br} \sum_{j=1}^b \sum_{k=1}^r Y_{ijk}, & \bar{Y}_{.j} &= \frac{1}{ar} \sum_{i=1}^a \sum_{k=1}^r Y_{ijk}\end{aligned}$$

Then,

$$\begin{aligned}\hat{\mu} &= \bar{Y}, & (\widehat{\tau\gamma})_{ij} &= \bar{Y}_{ij} - \hat{\mu} - \hat{\tau}_i - \hat{\gamma}_j \\ \hat{\tau}_i &= \bar{Y}_{i.} - \bar{Y}, & \hat{\gamma}_j &= \bar{Y}_{.j} - \bar{Y}\end{aligned}$$

and

$$\widehat{\sigma^2} = \frac{1}{N-1} \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (Y_{ijk} - \hat{\mu} - \hat{\tau}_i - \hat{\gamma}_j - (\widehat{\tau\gamma})_{ij})^2$$

Example: point estimation


Here's how to find estimates and confidence intervals in R

```
battery = read.csv("https://vicpena.github.io/doe/battery.csv")
battery$type = factor(battery$type)
battery$temp = factor(battery$temp)
options(contrasts = c("contr.sum", "contr.poly"))
mod = aov(life ~ type*temp, data = battery)
dummy.coef(mod)
confint(mod)
```

Not printing the results here for concreteness

Example: Interaction plot

```
emmip(mod, type ~ temp)
```



Global F tests

We have three F tests

- ▶ One for τ_i

$$H_{0,\tau} : \tau_i = 0, \text{ for all } i$$

$$H_{1,\tau} : \text{at least one } \tau_i \neq 0$$

- ▶ One for γ_j

$$H_{0,\gamma} : \gamma_j = 0, \text{ for all } j$$

$$H_{1,\gamma} : \text{at least one } \gamma_j \neq 0$$

- ▶ And another one for $(\tau\gamma)_{ij}$

$$H_{0,\tau\gamma} : (\tau\gamma)_{ij} = 0, \text{ for all } i, j$$

$$H_{1,\tau\gamma} : \text{at least one } (\tau\gamma)_{ij} \neq 0$$

Sums of squares

$$SS_{\text{total}} = SS_{\tau} + SS_{\gamma} + SS_{\tau\gamma} + SS_{\text{error}},$$

where

$$SS_{\text{total}} = \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (\bar{Y}_{.j} - \bar{Y})^2$$

$$SS_{\tau} = br \sum_{i=1}^a (\bar{Y}_{i.} - \bar{Y})^2$$

$$SS_{\gamma} = ar \sum_{j=1}^b (\bar{Y}_{.j} - \bar{Y})^2$$

$$SS_{\tau\gamma} = r \sum_{i=1}^a \sum_{j=1}^b (\bar{Y}_{ij} - \bar{Y}_{i.} - \bar{Y}_{.j} + \bar{Y})^2$$

$$SS_{\text{error}} = \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^r (Y_{ijk} - \bar{Y}_{ij})^2$$

ANOVA table

	df	SS	MS	F	p
τ	$a - 1$	SS_{τ}	$MS_{\tau} = \frac{SS_{\tau}}{a-1}$	$f_{\tau, \text{obs}} = \frac{MS_{\tau}}{MS_{\text{error}}}$	p_{τ}
γ	$b - 1$	SS_{γ}	$MS_{\gamma} = \frac{SS_{\gamma}}{b-1}$	$f_{\gamma, \text{obs}} = \frac{MS_{\gamma}}{MS_{\text{error}}}$	p_{γ}
$(\tau\gamma)$	$(a - 1)(b - 1)$	$SS_{\tau\gamma}$	$MS_{\tau\gamma} = \frac{SS_{\tau\gamma}}{(a-1)(b-1)}$	$f_{\tau\gamma, \text{obs}} = \frac{MS_{\tau\gamma}}{MS_{\text{error}}}$	$p_{\tau\gamma}$
Error	ν	SS_{error}	$MS_{\text{error}} = \frac{SS_{\text{error}}}{\nu}$		

Example: ANOVA table

```
summary(mod)
```


	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
## type	2	10728	5364	7.960	0.00192	**
## temp	2	39115	19557	29.020	1.88e-07	***
## type:temp	4	9670	2417	3.587	0.01803	*
## Residuals	27	18196	674			
## ---						
## Signif. codes:	0	'***'	0.001	'**'	0.01	'*' 0.05 '.' 0.1

Significant treatment effects and interaction at $\alpha = 0.05$.


Assumptions

- ▶ We still have assumptions of independence, normality, and equality of variances
- ▶ We can check them as we did with the blocked design
- ▶ Our model has an interaction: if it isn't significant, we can drop it from our model and fit an additive model

Example: checking equality of variances



Example: checking normality



Pairwise comparisons

- ▶ As before, global tests don't tell us where the differences are
- ▶ For that, we can still use TukeyHSD